

In [1]:

```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.utils import class_weight
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
from scipy.sparse import hstack

from tqdm import tqdm
import os
import datetime

import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.initializers import RandomNormal
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.layers import Input, Embedding, Dense, Flatten, concatenate
from tensorflow.keras.layers import BatchNormalization, Activation, CuDNNLSTM
from tensorflow.keras.models import Model, Sequential
```

```
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, Tensor
```

Splitting data into Train and Test

In [2]:

```
preprocessed_data = pd.read_csv('preprocessed_data.csv')
preprocessed_data.head(2)
```

Out[2]:

	Unnamed: 0	Unnamed: 0.1	id	teacher_id	teacher
0	0	8393	p205479	2bf07ba08945e5d8b2a3f269b2b3cfe5	
1	1	37728	p043609	3f60494c61921b3b43ab61bdde2904df	

2 rows × 21 columns

In [3]:

```
y = preprocessed_data['project_is_approved'].values
X = preprocessed_data.drop(['project_is_approved'], axis=1)
X.shape
```

Out[3]:

(109248, 20)

In [4]:

```
# train test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y)
X_test.shape
```

Out[4]:

(21850, 20)

In [5]:

```
X_train.head(2)
```

Out[5]:

	Unnamed: 0	Unnamed: 0.1	id	teacher_id	teacher_name
109078	109078	109837	p217509	c15f13b993b5478c32dcd13b9ebee016	
108069	108069	136104	p012508	6b93a7baa79ac6baade64f7c3484d970	

In [6]:

```
y_train
```

Out[6]:

array([1, 1, 1, ..., 1, 1, 1], dtype=int64)

In [7]:

```
X_test.head(2)
```

Out[7]:

	Unnamed: 0	Unnamed: 0.1	id	teacher_id	te
50340	50340	7583	p162251	49424684316a63797035055963389a0a	
65013	65013	32550	p204712	2b9f09cb8ce51fb624db6ba45f66d9ba	

In [8]:

```
y_test
```

Out[8]:

```
array([1, 1, 1, ..., 1, 0, 1], dtype=int64)
```

In [9]:

```
def loadGloveModel(gloveFile):  
    print ("Loading Glove Model")  
    f = open(gloveFile, 'r', encoding="utf8")  
    model = {}  
    for line in tqdm(f):  
        splitLine = line.split()  
        word = splitLine[0]  
        embedding = np.array([float(val) for val in splitLine[1:]])  
        model[word] = embedding  
    print ("Done.", len(model), " words loaded!")  
    return model  
model = loadGloveModel('glove.42B.300d.txt')
```

951it [00:00, 9438.94it/s]

Loading Glove Model

1917495it [03:28, 9206.60it/s]

Done. 1917495 words loaded!

In [10]:

```
t = Tokenizer()
t.fit_on_texts(X_train['essay'].values)
vocab_size = len(t.word_index) + 1
# integer encode the documents
X_train_encoded_docs = t.texts_to_sequences(X_train['essay'].values)
X_test_encoded_docs = t.texts_to_sequences(X_test['essay'].values)

# pad documents to a max length of 300 words
max_length = 300
X_train_padded_docs = pad_sequences(X_train_encoded_docs, maxlen=max_length,
X_test_padded_docs = pad_sequences(X_test_encoded_docs, maxlen=max_length, padding='post')

embedding_matrix = np.zeros((vocab_size, max_length))

for word, i in tqdm(t.word_index.items()):
    embedding_vector = model.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector

print(X_train_padded_docs.shape)
print(X_test_padded_docs.shape)
```

```
100%|███████████████████████████████████████████████████████████████████████|
██████████████████████████████████████████████████████████████████████████| 61833/61833 [00:00<00:00, 412605.36it/s]
```

(87398, 300)
(21850, 300)

In [11]:

```
print(vocab_size)
print(embedding_matrix.shape)
print(embedding_matrix[1])
```

```
61834
(61834, 300)
[-2.4837e-01 -4.5461e-01  3.9227e-02 -2.8422e-01 -3.1852e-02
 2.6355e-01
 -4.6323e+00  1.3890e-02 -5.3928e-01 -8.4454e-02  6.1556e-02 -
 4.1552e-01
 -1.4599e-01 -5.9321e-01 -2.8738e-02 -3.4991e-02 -2.9698e-01 -
 7.9850e-02
  2.7312e-01  2.2040e-01 -8.9859e-02  8.8265e-04 -4.1991e-01 -
 1.2536e-01
 -5.4629e-02  3.0550e-02  1.9340e-01 -6.3945e-02  2.7405e-02
 5.1193e-02
 -3.8656e-01 -1.1085e-01  1.7259e-01  2.9804e-01 -3.5183e-01
 1.3150e-01
 -5.4006e-01 -7.6677e-01 -5.5168e-04  1.3076e-01  2.5101e-02
 6.2106e-01
 -2.4797e-01 -3.9790e-01 -3.6116e-01 -5.1967e-01  3.0138e-02 -
 5.2436e-02
  6.9281e-02  3.5252e-02 -2.1402e-01  2.4836e-01 -1.5693e-01
 1.2829e-01
  3.5425e-01 -1.6080e-01 -5.0720e-03 -3.0656e-01 -2.9514e-01 -
 1.3554e-01
 -1.4385e-01 -4.0552e-01  5.7233e-01 -2.7670e-01  3.0519e-01
 1.5586e-01
  1.6086e-02 -2.2009e-01  4.8589e-01 -4.1384e-01  2.0546e-01
 4.0491e-01
  4.1558e-02 -1.3542e-01  2.2544e-01 -2.3629e-01  1.5193e-01 -
 1.0859e-02
 -8.2662e-02 -5.5484e-01 -6.1584e-02 -1.1112e-01 -1.1982e-01 -
 3.7064e-01
  1.6501e-01  4.4063e-01 -3.3883e-01 -5.7676e-01  5.0847e-01 -
 3.5707e-02
 -5.9233e-02  3.0748e-02 -2.7689e-01 -7.0433e-02  2.7786e-02 -
 5.9336e-01
 -2.8220e+00 -1.0052e-01  6.7168e-01 -1.7046e-01 -2.5902e-01
 2.7938e-01
  3.9992e-01  3.7480e-02 -2.6409e-01 -2.6378e-01  2.0645e-01
 1.7564e-01
 -8.0807e-02 -3.8376e-01  2.6602e-01  3.6214e-01 -9.5112e-02
 3.5199e-01
 -8.6994e-01 -1.5747e-01 -2.2550e-01 -6.4948e-02 -2.4845e-01
 1.5038e-01
 -3.2951e-01 -2.2285e-01 -2.5509e-02 -2.9725e-01 -3.7715e-01
 8.9296e-02
 -3.4399e-02  3.3640e-01  3.5534e-01  3.8253e-01  1.7646e-01
 1.3305e-01
```


-3.2743e-01	-4.7115e-01	2.4673e-01	-1.5964e-01	1.8212e-01	-
4.1241e-01					
9.8565e-02	3.8118e-01	3.3043e-01	5.1987e-02	-2.1824e-01	
2.2214e-01					
-5.9450e-02	-6.3743e-02	4.3723e-01	1.1068e-01	4.7444e-01	
5.6891e-01					
3.1123e-01	-2.0272e-01	8.0078e-02	-4.3905e-01	-1.2246e-01	-
2.5057e-02					
-5.7162e-02	1.4250e-01	9.4468e-02	1.2991e-01	1.0444e-01	-
3.9447e-01					
-2.9337e-01	-2.0466e-01	2.0815e-01	-1.6010e-01	-1.4665e-01	
5.4511e-01					
2.9740e-01	-2.2959e-01	-1.7050e-01	-6.2371e-02	-5.0399e-01	-
3.8000e-01					
-3.9528e-01	5.7552e-01	-4.6892e-01	-4.3308e-01	1.5018e-01	-
4.1179e-02					
6.2157e-01	1.9874e-02	-1.1969e-01	-2.5611e-01	2.6602e-01	-
3.7383e-01					
1.2936e-01	-5.0006e-02	-1.1554e-01	-1.7163e-01	-4.2430e-01	
1.9844e-01					
5.0611e-01	-1.1093e-01	-1.3939e-01	-5.9377e-01	6.7338e-01	
3.8497e-01					
6.2604e-01	-2.0128e-01	3.0058e-01	-1.3946e-01	-1.6186e-01	
1.2168e-01					
-1.8410e-02	6.1356e-01	-1.9887e-01	1.9250e-01	8.4372e-03	-
5.0757e-01					
3.5858e-01	-4.9729e-01	-4.4725e-01	2.1423e-02	-2.0769e-01	
8.3729e-02					
2.2032e-01	1.4404e-01	1.2590e-03	-4.4309e-01	-1.7242e-01	-
3.5300e-01					
-2.9477e-01	3.2898e-01	-3.1910e+00	3.8910e-01	3.5654e-01	
5.2134e-02					
2.0576e-01	-8.8649e-02	1.6398e-01	1.1203e-01	2.8590e-01	
2.8940e-01					
-4.4349e-01	9.1036e-01	-3.0902e-01	-1.3985e-01	-3.9499e-01	-
2.7299e-02					
-1.5201e-01	8.4418e-02	-3.7196e-01	4.9827e-02	1.4128e-01	-
1.5126e-01					
-1.6107e-01	4.0226e-03	1.6799e-01	-2.5429e-01	-1.5074e-01	-
5.7409e-01					
-1.5611e-01	6.8407e-02	2.4832e-01	1.6828e-01	7.2764e-02	-
8.6728e-02					
2.1982e-03	1.3593e-01	7.0224e-01	-4.5976e-01	-2.4506e-01	-
3.3874e-01					
-1.0952e-01	2.4698e-01	-5.5919e-01	-3.8866e-01	-1.3372e-01	
9.1943e-02					
-1.0543e-01	-3.1319e-01	-2.9952e-01	-2.0611e-01	1.7976e-01	
4.5800e-01					
-7.2402e-02	1.6118e-01	-4.1649e-01	-3.0103e-01	2.3234e-01	-
5.0139e-02					
1.0026e-01	3.8974e-01	-6.1342e-02	2.6626e-01	-1.5671e-01	
7.5136e-02					
-4.2926e-01	-1.2025e-01	8.2736e-02	-6.2469e-01	4.4267e-02	

```
6.0673e-01  
-1.2458e-01 -1.5443e-01 -1.6339e-01 5.3097e-02 1.5458e-01 -  
3.8053e-01]
```

1.4 Encoding Categorical and Numerical features

1.4.1 encoding categorical features: clean_categories

In [12]:

```
vectorizer_cat = CountVectorizer()  
vectorizer_cat.fit(X_train['clean_categories'].values) # fit has to happen on  
  
X_train_cc_ohe = vectorizer_cat.transform(X_train['clean_categories'].values)  
X_test_cc_ohe = vectorizer_cat.transform(X_test['clean_categories'].values)  
  
print("After vectorizations")  
print(X_train_cc_ohe.shape, y_train.shape)  
print(X_test_cc_ohe.shape, y_test.shape)  
print(vectorizer_cat.get_feature_names())
```

```
After vectorizations  
(87398, 9) (87398,)  
(21850, 9) (21850,)  
['appliedlearning', 'care_hunger', 'health_sports', 'history_c  
ivics', 'literacy_language', 'math_science', 'music_arts', 'sp  
ecialneeds', 'warmth']
```

1.4.2 encoding categorical features: clean_subcategories

In [13]:

```
vectorizer_subcat = CountVectorizer()
vectorizer_subcat.fit(X_train['clean_subcategories'].values) # fit has to happen on training data

X_train_csc_ohe = vectorizer_subcat.transform(X_train['clean_subcategories']).toarray()
X_test_csc_ohe = vectorizer_subcat.transform(X_test['clean_subcategories']).toarray()

print("After vectorizations")
print(X_train_csc_ohe.shape, y_train.shape)
print(X_test_csc_ohe.shape, y_test.shape)
print(vectorizer_subcat.get_feature_names())
```

```
After vectorizations
(87398, 30) (87398,)
(21850, 30) (21850,)
['appliedsciences', 'care_hunger', 'charactereducation', 'civics_government', 'college_careerprep', 'communityservice', 'earlydevelopment', 'economics', 'environmentalscience', 'esl', 'extracurricular', 'financialliteracy', 'foreignlanguages', 'gym_fitness', 'health_lifescience', 'health_wellness', 'history_geography', 'literacy', 'literature_writing', 'mathematics', 'music', 'nutritioneducation', 'other', 'parentinvolvement', 'performingarts', 'socialsciences', 'specialneeds', 'teamsports', 'visualarts', 'warmth']
```

1.4.3 encoding categorical features: school_state

In [14]:

```
vectorizer_school_state = CountVectorizer()
vectorizer_school_state.fit(X_train['school_state'].values)

X_train_state_ohe = vectorizer_school_state.transform(X_train['school_state'].values)
X_test_state_ohe = vectorizer_school_state.transform(X_test['school_state'].values)

print("After vectorizations")
print(X_train_state_ohe.shape, y_train.shape)
print(X_test_state_ohe.shape, y_test.shape)
print(vectorizer_school_state.get_feature_names())
```

```
After vectorizations
(87398, 51) (87398,)
(21850, 51) (21850,)
['ak', 'al', 'ar', 'az', 'ca', 'co', 'ct', 'dc', 'de', 'fl',
'ga', 'hi', 'ia', 'id', 'il', 'in', 'ks', 'ky', 'la', 'ma', 'm
d', 'me', 'mi', 'mn', 'mo', 'ms', 'mt', 'nc', 'nd', 'ne', 'n
h', 'nj', 'nm', 'nv', 'ny', 'oh', 'ok', 'or', 'pa', 'ri', 's
c', 'sd', 'tn', 'tx', 'ut', 'va', 'vt', 'wa', 'wi', 'wv', 'w
y']
```

1.4.4 encoding categorical features: teacher_prefix

In [15]:

```
vectorizer_prefix = CountVectorizer()
vectorizer_prefix.fit(X_train['teacher_prefix'].values)

X_train_teacher_ohe = vectorizer_prefix.transform(X_train['teacher_prefix'].values)
X_test_teacher_ohe = vectorizer_prefix.transform(X_test['teacher_prefix'].values)

print("After vectorizations")
print(X_train_teacher_ohe.shape, y_train.shape)
print(X_test_teacher_ohe.shape, y_test.shape)
print(vectorizer_prefix.get_feature_names())
```

```
After vectorizations
(87398, 5) (87398,)
(21850, 5) (21850,)
['dr', 'mr', 'mrs', 'ms', 'teacher']
```

1.4.5 encoding categorical features: project_grade_category

In [16]:

```
vectorizer_grade = CountVectorizer()
vectorizer_grade.fit(X_train['project_grade_category'].values)

X_train_grade_ohe = vectorizer_grade.transform(X_train['project_grade_category'])
X_test_grade_ohe = vectorizer_grade.transform(X_test['project_grade_category'])

print("After vectorizations")
print(X_train_grade_ohe.shape, y_train.shape)
print(X_test_grade_ohe.shape, y_test.shape)
print(vectorizer_grade.get_feature_names())
```

```
After vectorizations
(87398, 4) (87398,)
(21850, 4) (21850,)
['grades_3_5', 'grades_6_8', 'grades_9_12', 'grades_prek_2']
```

1.4.6 encoding numerical features: price

In [17]:

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
scaler.fit(X_train['price'].values.reshape(-1,1))
X_train_price=scaler.transform(X_train['price'].values.reshape(-1,1))
X_test_price=scaler.transform(X_test['price'].values.reshape(-1,1))

print(X_train_price.shape)
print(X_test_price.shape)
print(X_train_price)
print(X_test_price)
```

```
(87398, 1)
(21850, 1)
[[0.02126153]
 [0.00848041]
 [0.02222569]
 ...
 [0.02341289]
 [0.02103149]
 [0.04478043]]
[[0.04994129]
 [0.01193498]
 [0.03926652]
 ...
 [0.01223203]
 [0.04165792]
 [0.01909417]]
```

1.4.7 encoding numerical features: teacher_number_of_previously_posted_projects

In [18]:

```
scaler = MinMaxScaler()
scaler.fit(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
X_train_pre_proj=scaler.transform(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
X_test_pre_proj=scaler.transform(X_test['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))

print(X_train_pre_proj.shape)
print(X_test_pre_proj.shape)
```

```
(87398, 1)
(21850, 1)
```

1.4.8 encoding numerical features: quantity

In [19]:

```
scaler = MinMaxScaler()
scaler.fit(X_train['quantity'].values.reshape(-1,1))
X_train_quantity=scaler.transform(X_train['quantity'].values.reshape(-1,1))
X_test_quantity=scaler.transform(X_test['quantity'].values.reshape(-1,1))

print(X_train_quantity.shape)
print(X_test_quantity.shape)
```

(87398, 1)

(21850, 1)

In []:

In [20]:

```
X_train_cn = hstack((X_train_cc_ohc,X_train_csc_ohc,X_train_teacher_ohc,X_train_gr
                    X_train_state_ohc,X_train_price,X_train_quantity,X_train_t
X_test_cn = hstack((X_test_cc_ohc,X_test_csc_ohc,X_test_teacher_ohc,X_test_gr
                    X_test_state_ohc,X_test_price,X_test_quantity,X_test_p
```

In [21]:

```
print(X_train_cn.shape)
print(X_test_cn.shape)
```

(87398, 102)

(21850, 102)

In [22]:

```
print('vocabulary size is ',vocab_size)
```

vocabulary size is 61834

In [1]:

```
'''
These cells are to load the data if the system crashes during training
'''

import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.utils import class_weight
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
from scipy.sparse import hstack

from tqdm import tqdm
import os
import datetime

import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.initializers import RandomNormal
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.layers import Input, Embedding, Dense, Flatten, concatenate
```



```
from tensorflow.keras.layers import BatchNormalization, Activation, CuDNNLSTM
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, TensorBoard
vocab_size = 61834
```

In [23]:

#Saving processed data

```
from sklearn.externals import joblib

joblib.dump(embedding_matrix, 'embedding_matrix3.pkl')
joblib.dump(X_train_padded_docs, 'X_train_padded_docs3.pkl')
joblib.dump(X_train_cn, 'X_train_cn3.pkl')
joblib.dump(y_train, 'y_train3.pkl')

joblib.dump(X_test_padded_docs, 'X_test_padded_docs3.pkl')
joblib.dump(X_test_cn, 'X_test_cn3.pkl')
joblib.dump(y_test, 'y_test3.pkl')
```

Out[23]:

```
['y_test3.pkl']
```

In [2]:

#Loading processed data

```
from sklearn.externals import joblib

embedding_matrix = joblib.load('embedding_matrix3.pkl')
X_train_padded_docs = joblib.load('X_train_padded_docs3.pkl')
X_train_cn = joblib.load('X_train_cn3.pkl')
y_train = joblib.load('y_train3.pkl')

X_test_padded_docs = joblib.load('X_test_padded_docs3.pkl')
X_test_cn = joblib.load('X_test_cn3.pkl')
y_test = joblib.load('y_test3.pkl')
```

In [3]:

```
from sklearn.metrics import roc_auc_score
def auc( y_true, y_pred ) :
    score = tf.py_func( lambda y_true, y_pred : roc_auc_score( y_true, y_pred
                                                                [y_true, y_pred],
                                                                'float32',
                                                                stateful=True,
                                                                name='sklearnAUC' )
                        y_true, y_pred )
    return score
```

Model's architecture

In [4]:

```
# Essay Layers
essay_input = Input(shape=(len(X_train_padded_docs[0]),), name='essay_input')

essay_input_1 = Embedding(input_dim=vocab_size, output_dim=300, input_length=1,
                           weights=[embedding_matrix], trainable=False)(essay_input)

essay_input_1 = CuDNNLSTM(units = 64,
                           kernel_initializer= 'he_normal',
                           return_sequences=True)(essay_input_1)
essay_input_1 = Flatten()(essay_input_1)

# Categorical and numerical Layers

cat_num_input = Input(shape=(X_train_cn.shape[1],), name='cat_num_input')

cn_input = Embedding(input_dim=X_train_cn.shape[1], output_dim=128, input_length=1)(cat_num_input)

conv1d_1 = Conv1D(filters=64, kernel_size=(5), padding='same')(cn_input)

flat_2 = Flatten()(conv1d_1)

# Concatinating all the above Layers
x = concatenate([essay_input_1, flat_2])

# Dense Layers
output = Dense(units = 624,
                activation = 'relu',
                kernel_initializer = 'he_normal')(x)
output = Dropout(0.7)(output)
output = BatchNormalization()(output)

output = Dropout(0.7)(output)
output = BatchNormalization()(output)
output = Dense(units = 512,
                activation = 'relu',
                kernel_initializer = 'he_normal')(output)

output = Dropout(0.8)(output)
output = BatchNormalization()(output)
output = Dense(units = 512,
                activation = 'relu',
                kernel_initializer = 'he_normal')(output)
output = Dropout(0.45)(output)
output = Dense(1, activation='sigmoid', name='output')(output)

model_3 = Model(inputs = [essay_input, cat_num_input],
                 outputs=output)

model_3.compile(loss = 'binary_crossentropy',
```

```
optimizer = 'adam',  
metrics = ['accuracy', auc])
```

WARNING:tensorflow:From c:\users\addu\appdata\local\programs\python\python37\lib\site-packages\tensorflow\python\keras\initializers.py:119: calling RandomUniform.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From c:\users\addu\appdata\local\programs\python\python37\lib\site-packages\tensorflow\python\ops\init_ops.py:1251: calling VarianceScaling.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

WARNING:tensorflow:Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

WARNING:tensorflow:Large dropout rate: 0.8 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

WARNING:tensorflow:From <ipython-input-3-85586a596984>:7: py_func (from tensorflow.python.ops.script_ops) is deprecated and will be removed in a future version.

Instructions for updating:

tf.py_func is deprecated in TF V2. Instead, there are two options available in V2.

- tf.py_function takes a python function which manipulates tf eager

tensors instead of numpy arrays. It's easy to convert a tf eager tensor to

an ndarray (just call tensor.numpy()) but having access to eager tensors

means `tf.py_function`s can use accelerators such as GPUs as well as

being differentiable using a gradient tape.

- tf.numpy_function maintains the semantics of the deprecated tf.py_func

(it is not differentiable, and manipulates numpy arrays).

It drops the

stateful argument making all functions stateful.

WARNING:tensorflow:From c:\users\addu\appdata\local\programs\python\python37\lib\site-packages\tensorflow\python\ops\nn_impl.py:180: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in

a future version.

Instructions for updating:

Use `tf.where` in 2.0, which has the same broadcast rule as `np.where`

Training the model with train data

In [5]:

```
log_dir="logs\\model_3_log" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard = TensorBoard(log_dir=log_dir)

mcp_save = ModelCheckpoint('.model3_best_weights.hdf5', save_best_only=True,

callbacks_list = [mcp_save, tensorboard]
```

In [6]:

```
history = model_3.fit([X_train_padded_docs, X_train_cn.todense()],
                      y_train,
                      batch_size = 1024,
                      epochs = 50,
                      callbacks=callbacks_list,
                      validation_split=0.2,
                      verbose = 2)
```

Train on 69918 samples, validate on 17480 samples

Epoch 1/50

69918/69918 - 8s - loss: 0.5265 - acc: 0.8183 - auc: 0.5096
- val_loss: 0.5312 - val_acc: 0.8478 - val_auc: 0.5765

Epoch 2/50

69918/69918 - 5s - loss: 0.4626 - acc: 0.8418 - auc: 0.5328
- val_loss: 0.5188 - val_acc: 0.8478 - val_auc: 0.5712

Epoch 3/50

69918/69918 - 5s - loss: 0.4460 - acc: 0.8478 - auc: 0.5148
- val_loss: 0.5211 - val_acc: 0.8478 - val_auc: 0.4286

Epoch 4/50

69918/69918 - 5s - loss: 0.4321 - acc: 0.8488 - auc: 0.5332
- val_loss: 0.5042 - val_acc: 0.8478 - val_auc: 0.4256

Epoch 5/50

69918/69918 - 5s - loss: 0.4249 - acc: 0.8488 - auc: 0.5603
- val_loss: 0.5722 - val_acc: 0.8478 - val_auc: 0.4244

Epoch 6/50

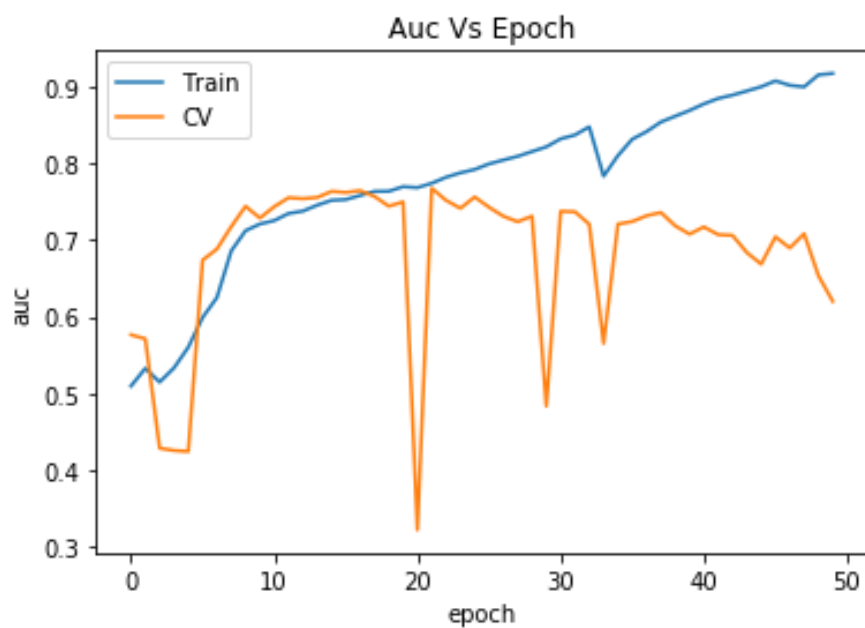
69918/69918 - 5s - loss: 0.4179 - acc: 0.8488 - auc: 0.5991
- val_loss: 0.5706 - val_acc: 0.8478 - val_auc: 0.6740

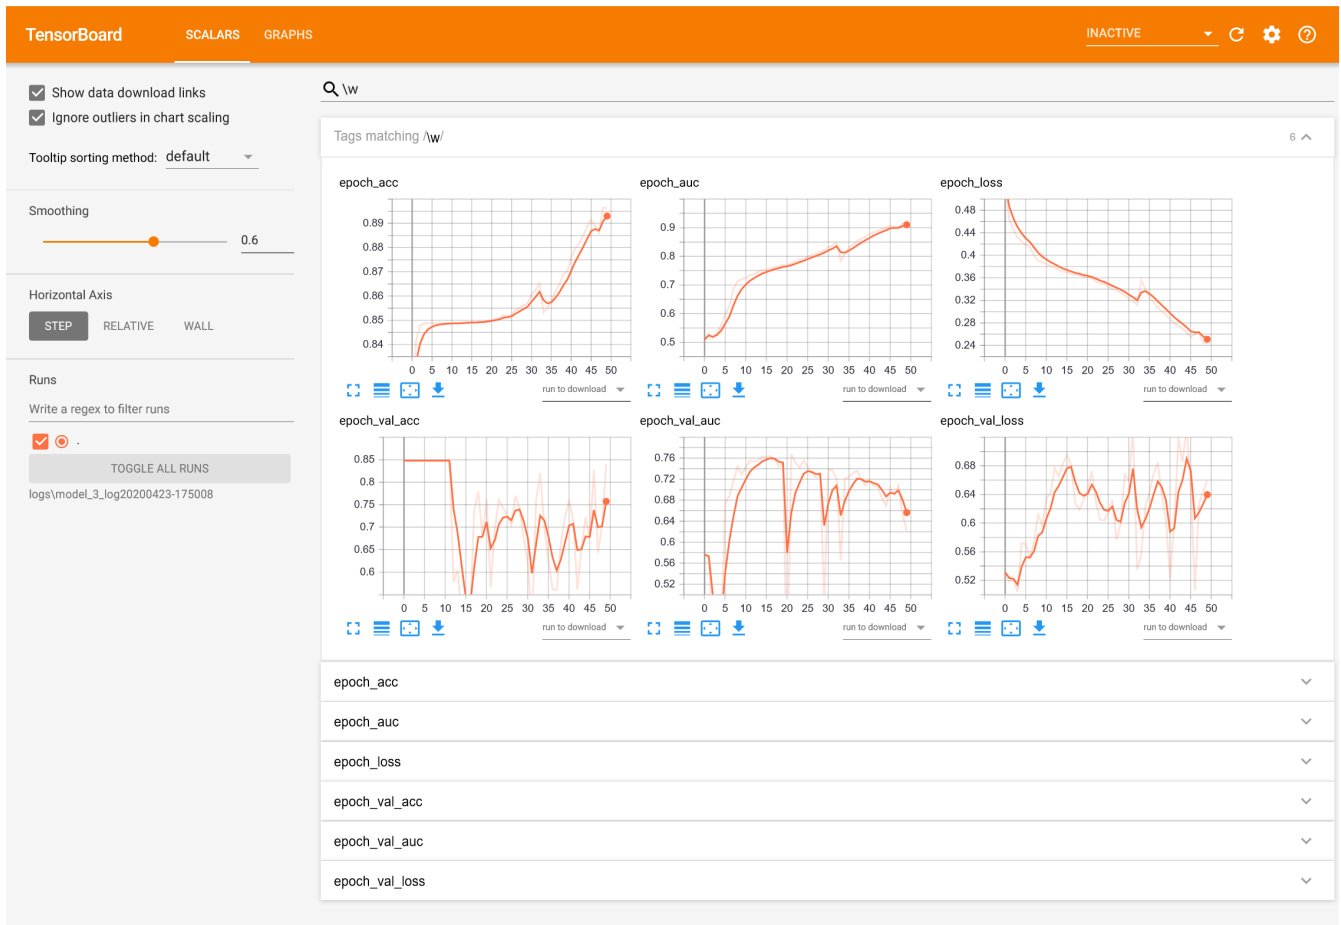
Epoch 7/50

Auc vs Epoch

In [7]:

```
# summarize history for accuracy
plt.plot(history.history['auc'])
plt.plot(history.history['val_auc'])
plt.title('Auc Vs Epoch')
plt.ylabel('auc')
plt.xlabel('epoch')
plt.legend(['Train', 'CV'], loc='right bottom')
plt.show()
```





Testing the model with test data

In [8]:

```
# Essay Layers
essay_input = Input(shape=(len(X_train_padded_docs[0]),), name='essay_input')

essay_input_1 = Embedding(input_dim=vocab_size,output_dim=300, input_length=1,
                           weights=[embedding_matrix], trainable=False)(essay_input)

essay_input_1 = CuDNNLSTM(units = 64,
                           kernel_initializer= 'he_normal',
                           return_sequences=True)(essay_input_1)
essay_input_1 = Flatten()(essay_input_1)

# Categorical and numerical Layers

cat_num_input = Input(shape=(X_train_cn.shape[1],), name='cat_num_input')

cn_input = Embedding(input_dim=X_train_cn.shape[1], output_dim=128, input_length=1)(cat_num_input)

conv1d_1 = Conv1D(filters=64, kernel_size=(5), padding='same')(cn_input)

flat_2 = Flatten()(conv1d_1)

# Concatinating all the above Layers
x = concatenate([essay_input_1, flat_2])

# Dense Layers
output = Dense(units = 624,
                activation = 'relu',
                kernel_initializer = 'he_normal')(x)
output = Dropout(0.7)(output)
output = BatchNormalization()(output)

output = Dropout(0.7)(output)
output = BatchNormalization()(output)
output = Dense(units = 512,
                activation = 'relu',
                kernel_initializer = 'he_normal')(output)

output = Dropout(0.8)(output)
output = BatchNormalization()(output)
output = Dense(units = 512,
                activation = 'relu',
                kernel_initializer = 'he_normal')(output)
output = Dropout(0.45)(output)
output = Dense(1, activation='sigmoid', name='output')(output)

model_3 = Model(inputs = [essay_input, cat_num_input],
                outputs=output)

model_3.compile(loss = 'binary_crossentropy',
```



```
optimizer = 'adam',  
metrics = ['accuracy', auc])
```

WARNING:tensorflow:Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

WARNING:tensorflow:Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

In [9]:

```
#loading the best weights obtained from training model on train data  
model_3.load_weights(".model3_best_weights.hdf5")
```

In [10]:

```
score=model_3.evaluate([X_test_padded_docs, X_test_cn.todense()],y_test,  
                        batch_size=1024, verbose=0)[2]
```

Auc score of model on unseen Test data

In [11]:

```
print('Test Auc obtained is ',score)
```

Test Auc obtained is 0.75855786

In [12]:

```
model_3.summary()
```

Model: "model_1"

Layer (type) Connected to	Output Shape	Param #
=====		
essay_input (InputLayer)	[(None, 300)]	0
=====		
cat_num_input (InputLayer)	[(None, 102)]	0
=====		
embedding_2 (Embedding) essay_input[0][0]	(None, 300, 300)	18550200
=====		
embedding_3 (Embedding) cat_num_input[0][0]	(None, 102, 128)	13056
=====		
cu_dnnlstm_1 (CuDNNLSTM) embedding_2[0][0]	(None, 300, 64)	93696
=====		
conv1d_1 (Conv1D) embedding_3[0][0]	(None, 102, 64)	41024
=====		
flatten_2 (Flatten) cu_dnnlstm_1[0][0]	(None, 19200)	0
=====		
flatten_3 (Flatten) conv1d_1[0][0]	(None, 6528)	0
=====		
concatenate_1 (Concatenate) flatten_2[0][0] flatten_3[0][0]	(None, 25728)	0
=====		
dense_3 (Dense) concatenate_1[0][0]	(None, 624)	16054896
=====		
dropout_4 (Dropout)	(None, 624)	0

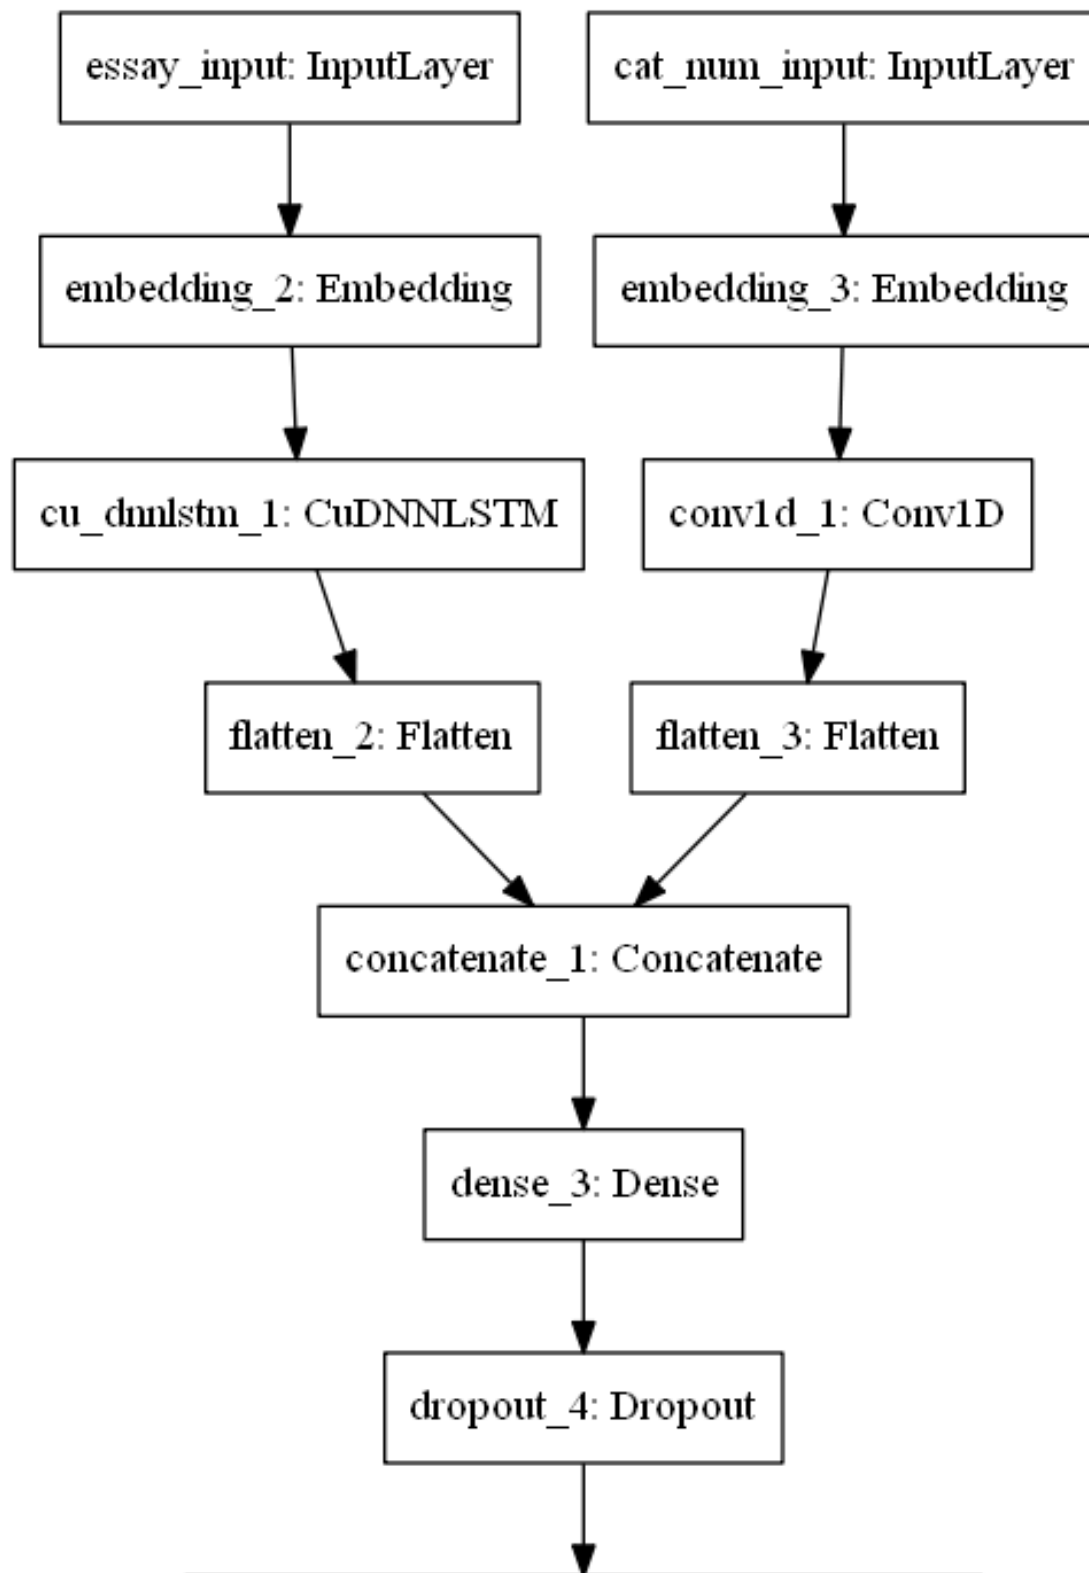
dense_3[0][0]		
batch_normalization_3 (BatchNormalizer) (None, 624)		2496
dropout_4[0][0]		
dropout_5 (Dropout) (None, 624)		0
batch_normalization_3[0][0]		
batch_normalization_4 (BatchNormalizer) (None, 624)		2496
dropout_5[0][0]		
dense_4 (Dense) (None, 512)		320000
batch_normalization_4[0][0]		
dropout_6 (Dropout) (None, 512)		0
dense_4[0][0]		
batch_normalization_5 (BatchNormalizer) (None, 512)		2048
dropout_6[0][0]		
dense_5 (Dense) (None, 512)		262656
batch_normalization_5[0][0]		
dropout_7 (Dropout) (None, 512)		0
dense_5[0][0]		
output (Dense) (None, 1)		513
dropout_7[0][0]		
=====		
=====		
Total params: 35,343,081		
Trainable params: 16,789,361		
Non-trainable params: 18,553,720		

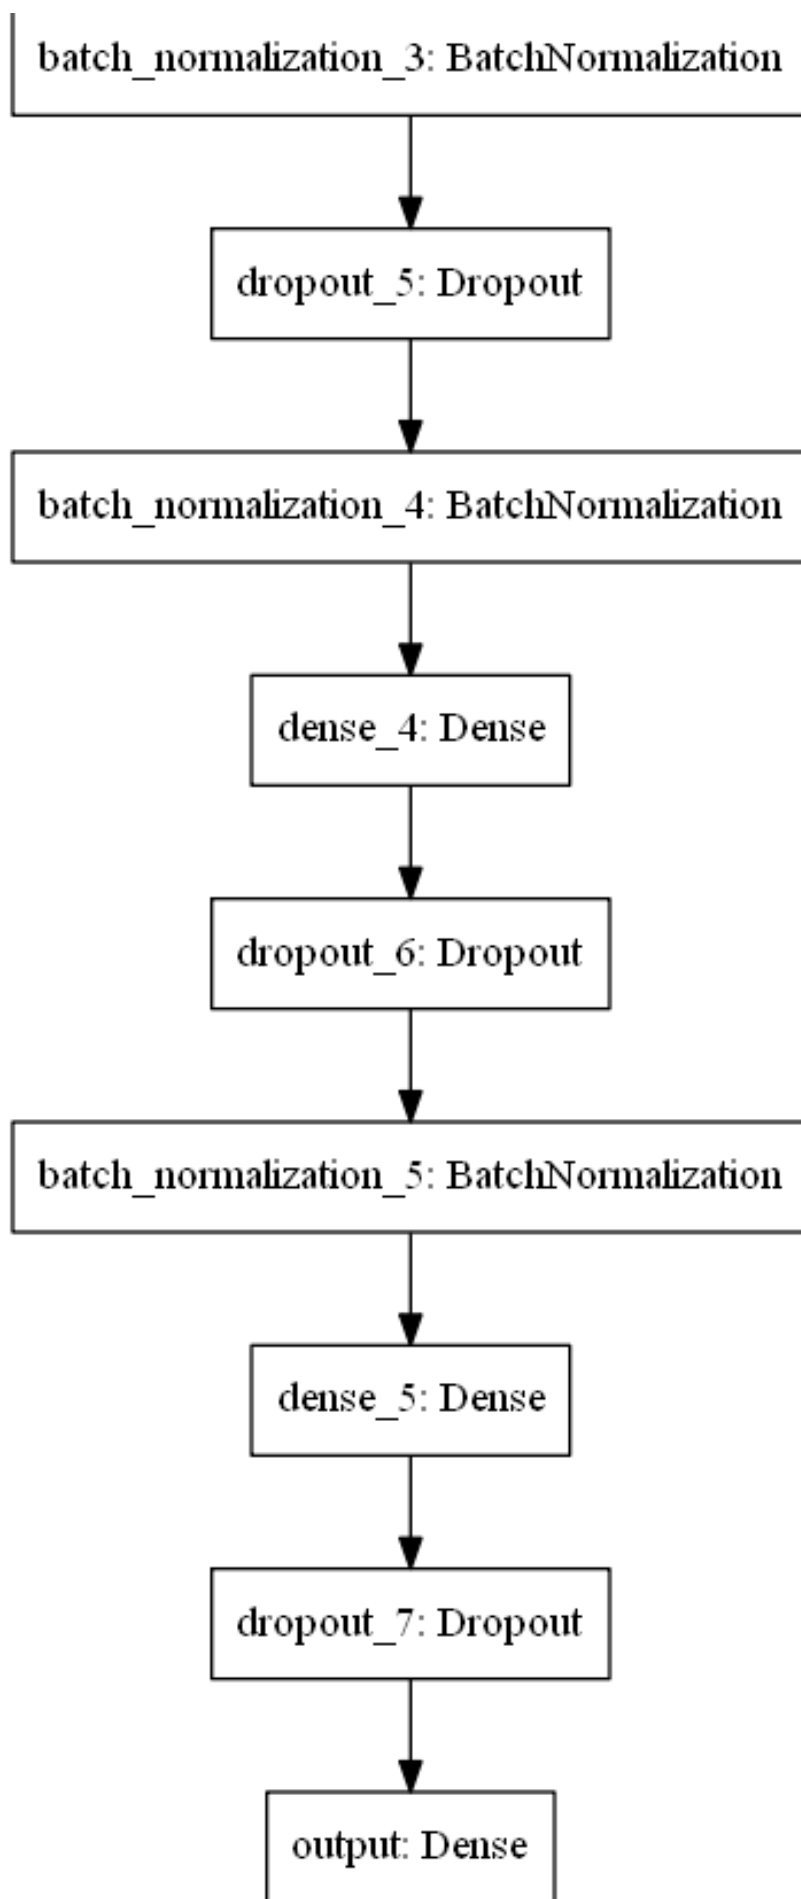
Model's Final Architecture

In [13]:

```
#drawing models
tf.keras.utils.plot_model(
    model_3,
    show_shapes=False,
    show_layer_names=True,
    rankdir='TB'
)
```

Out[13]:





Summary

- Created tensorflow model using text, categorical and numerical layers.
- Performed hyper parameter tuning manually on number of layers, activation functions and optimizers.
- Trained the network using the best obtained hyper parameters.
- Tested the model using Test data and obtained test aucroc score of 0.7585
- Printed the summary of the model along with it's image.